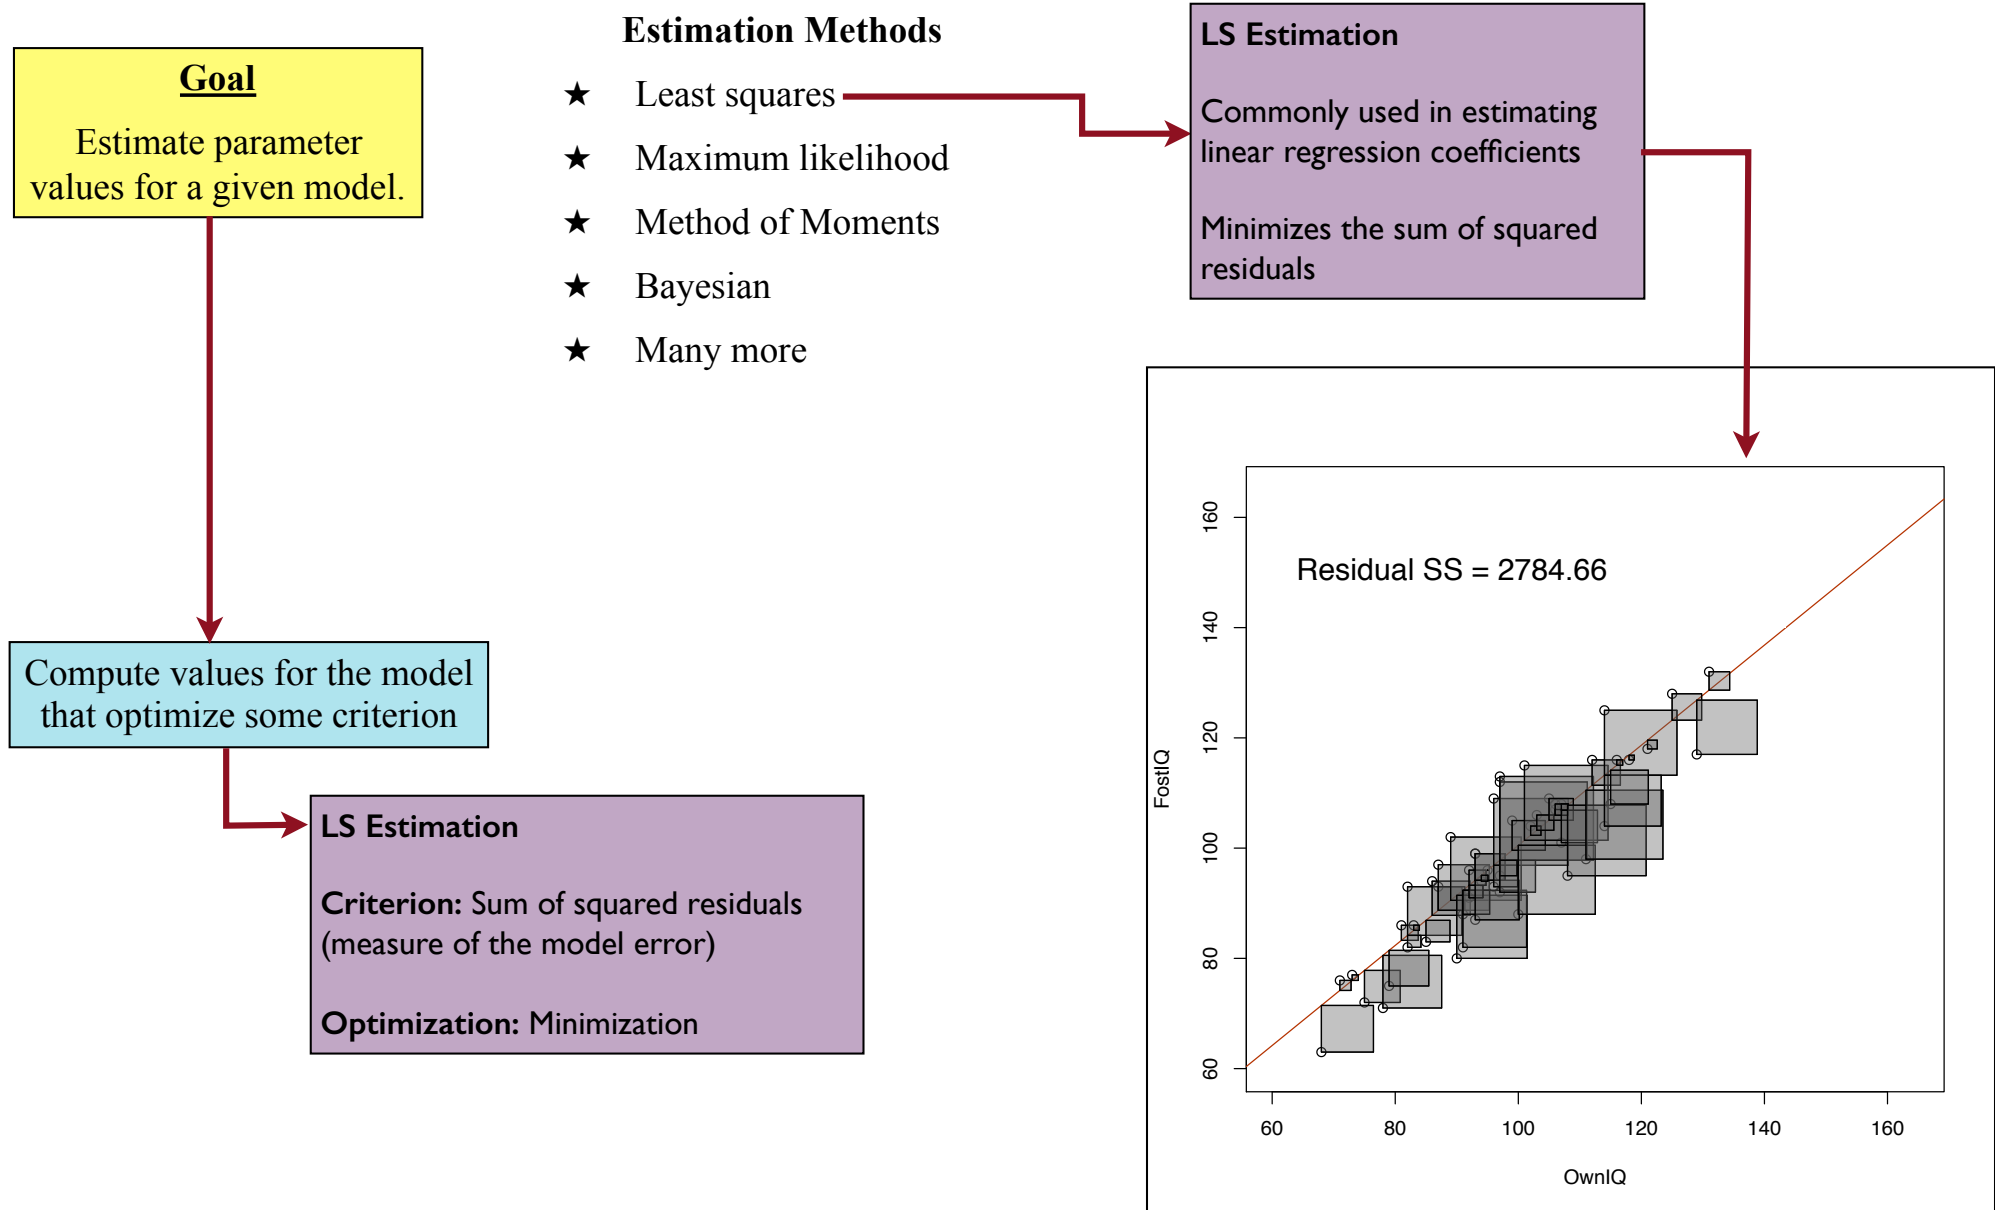


# **Maximum Likelihood Estimation**

# Estimation



# Maximum Likelihood

## Key Question

Given the data and model, what are the most likely value of the parameters?

$$L(x_1, x_2, x_3, \dots, x_n; \theta)$$

Likelihood function represents the **joint probability** or **likelihood** of observing the data that have been collected

Maximum likelihood (ML) provides framework for answering this question

and

provides results which have attractive properties favorable for inference


**Criterion:** likelihood function (or log-likelihood function or deviance function)

**Optimization:** Minimization

## Advantages of ML

- ★ ML yields a **global fit index** for the model on top of the parameter estimates
    - ✓ Can be used to compare models
  - ★ Estimators produced under ML have **desirable large-sample (asymptotic) properties**
    - ✓ Consistent, asymptotically normally distributed
    - ✓ With small samples this may not hold
- 
- ★ ML estimates are **always** approximate
  - ★ Approximations imply researchers **should not** get too hung up on rigid rules of practice
    - ✓ Use of 2(SE) vs 1.96(SE)
    - ✓ Inflexible cutoffs (0.05) are not warranted
  - ★ Notes only consider **regular problems**
    - ✓ ML solutions are possible
    - ✓ Assumption of sample data coming from hypothetical, infinitely sized population (repeated sampling scenario)

# Example 1 : Estimate the Mean



**The Model**

$$y_i = \beta_0 + \epsilon_i$$

The errors are normally distributed

Given the **data** and the **model**, what is the most likely value of the mean?

**Data:** The data are given

**Model:** Specified by the researcher

e.g., the marginal mean model is an LM with distributional assumption on the errors

**The Data**

In general are  $y_1, y_2, y_3, \dots, y_n$

**Distributional assumption is normal**

The probability density of  $y_i$  is  $f(y_i) = \frac{1}{\sqrt{2\pi\sigma^2}} \times \exp \left[ -\frac{(y_i - \mu)^2}{2\sigma^2} \right]$

With the further usual assumption that  $\mu_\epsilon = 0$  this simplifies to

$$f(y_i) = \frac{1}{\sqrt{2\pi\sigma^2}} \times \exp \left[ -\frac{y_i^2}{2\sigma^2} \right]$$

Probability model associates different probability densities with different  $y_i$ 's

Suppose  $\mu = 0$ ,  $\sigma^2 = 1$  and  $y_i = -1$  then the probability density is given by  $f(y_i) = f(-1)$


$$f(-1) = \frac{1}{\sqrt{2 \cdot \pi \cdot 1}} \times \exp \left[ -\frac{(-1)^2}{2 \cdot 1} \right] \\ = 0.2419707$$

The `dnorm()` function computes probability densities from a normal distribution in R.

```
> dnorm(-1, mean = 0, sd = sqrt(1) )
```

```
[1] 0.2419707
```

Note it takes the SD rather than the variance of the distribution



# Joint Probability Density

The **joint probability density** is the probability of observing  $y_1$  *and*  $y_2$  *and*  $y_3$  *and* ... *and*  $y_n$

Assumption of **independence** allows us to multiply each probability density function to obtain **joint density**

$$p(y_1, y_2, \dots, y_n) = f(y_1) \times f(y_2) \times \dots \times f(y_N)$$

$$p(y_1, y_2, \dots, y_n) = \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \times \exp\left(-\frac{y_1^2}{2\sigma^2}\right) \right] \times \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \times \exp\left(-\frac{y_2^2}{2\sigma^2}\right) \right] \times \dots \times \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \times \exp\left(-\frac{y_n^2}{2\sigma^2}\right) \right]$$

$$p(y_1, y_2, \dots, y_n) = \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right)^N \times \exp\left[-\frac{\sum y_i^2}{2\sigma^2}\right]$$

**joint density function**

---

In working with statistical models, typically the distributional assumption is on the errors. This whole process is carried out using the errors ( $\epsilon_i$ ) rather than the  $y$ 's

$$p(\epsilon_1, \epsilon_2, \dots, \epsilon_n) = \left( \frac{1}{\sqrt{2\pi\sigma_\epsilon^2}} \right)^N \times \exp\left[-\frac{\sum \epsilon_i^2}{2\sigma_\epsilon^2}\right]$$

**joint density of the errors**

# Likelihood

Then we have the joint probability density of the errors...

...and a model

$$p(\epsilon_1, \epsilon_2, \dots, \epsilon_n) = \left( \frac{1}{\sqrt{2\pi\sigma_\epsilon^2}} \right)^N \times \exp \left[ -\frac{\sum \epsilon_i^2}{2\sigma_\epsilon^2} \right]$$

$$y_i = \beta_0 + \epsilon_i$$

We can re-express  $\epsilon_i$  as  $y_i - \beta_0$  and substitute into the function for the joint density

$$p(\epsilon_1, \epsilon_2, \dots, \epsilon_n) = \left( \frac{1}{\sqrt{2\pi\sigma_\epsilon^2}} \right)^N \times \exp \left[ -\frac{\sum (y_i - \beta_0)^2}{2\sigma_\epsilon^2} \right]$$

---

Once we express the probability density as a function of parameters it is called the **likelihood function**.

$$\mathcal{L}(\beta; \mathbf{y}) = \left( \frac{1}{\sqrt{2\pi\sigma_\epsilon^2}} \right)^N \times \exp \left[ -\frac{\sum (y_i - \beta_0)^2}{2\sigma_\epsilon^2} \right]$$

**Likelihood function**

**Likelihood of the vector of betas given the data**

# Log-Likelihood

It is often easier to work mathematically with the log of the likelihood function (**log-likelihood**).

$$\log \mathcal{L}(\beta; \mathbf{y}) = \log \left( \left( \frac{1}{\sqrt{2\pi\sigma_\epsilon^2}} \right)^N \times \exp \left[ -\frac{\sum (y_i - \beta_0)^2}{2\sigma_\epsilon^2} \right] \right)$$

Choosing the natural log  
this is re-expressed as

$$\ell(\beta; \mathbf{y}) = -\frac{N}{2} \times \ln(2\pi\sigma_\epsilon^2) - \frac{1}{2\sigma_\epsilon^2} \times \sum (y_i - \beta_0)^2$$

Log-likelihood function

Do the math to convince yourself it reduces to this

- ★ Assume that the parameters of the error variance is known, and that the **only unknown** parameter to be estimated is the intercept

$$\sigma_\epsilon^2 = 15$$

- ★ Given the model and the data, object is to choose "best" value for the intercept
- ★ "Best" value here means that after substituting values in to the deviance function equation, the smallest deviance possible has been obtained (i.e., deviance is minimized)
- ★ In ML theory, this is the estimate (value) of  $\beta_0$  that is **maximally likely** given the data and the model—hence maximum likelihood



Everything in the log-likelihood function is known, except for  $\beta_0$

$$\ell(\beta; \mathbf{y}) = -\frac{N}{2} \times \ln(2\pi\sigma_\epsilon^2) - \frac{1}{2\sigma_\epsilon^2} \times \sum (y_i - \beta_0)^2$$

$N = 5$                        $\sigma_\epsilon^2 = 15$                        $y_i = \{12, 16.26, 8.50, 8, 13.65\}$

Substitute in the values for  $\sigma_\epsilon^2$ ,  $y_i$ , and  $N$


$$\ell(\beta; \mathbf{y}) = -\frac{5}{2} \times \ln(2 \times \pi \times 15) - \frac{1}{2 \times 15} \times [(12 - \beta_0)^2 + (16.26 - \beta_0)^2 + (8.50 - \beta_0)^2 + (8 - \beta_0)^2 + (13.65 - \beta_0)^2]$$

Given we want the largest (maximum) log-likelihood what value should we choose for  $\beta_0$ ?

- ★ Pick several candidate values for  $\beta_0$
- ★ Substitute them into the equation
- ★ Solve for deviance
- ★ Choose the value for  $\beta_0$  that produces the smallest deviance

For example, consider candidate values for  $\beta_0$  between 11 and 12

$\beta_0$	Log-likelihood
11	-13.06282
11.01	-13.06056
11.02	-13.05834
11.03	-13.05615
$\vdots$	$\vdots$



```
## Create a function to compute the log-likelihood
> LL = function(b0) {
  -5/2 * log(2 * pi * 15) - 1/(2 * 15) * ( (12 - b0)^2 + (8 - b0)^2 + (16.26 - b0)^2 +
    (13.65 - b0)^2 + (8.5 - b0)^2 )
}

## Try function
> LL(11)

[1] -13.06282
```

# Carry Out a Grid Search in R

```
> library(dplyr)
> library(ggplot2)

## Generate candidate values for b0
> candidates = data.frame(
  b0 = seq(from = 11, to = 12, by = 0.01)
)

## Generate the deviance values and store the results
> gridSearch = candidates %>%
  rowwise() %>%
  mutate(ll = LL(b0))

## Examine search grid
> head(gridSearch) +
```

Source: local data frame [6 x 2]

	b0	ll
	(dbl)	(dbl)
1	11.00	-13.06282
2	11.01	-13.06056
3	11.02	-13.05834
4	11.03	-13.05615
5	11.04	-13.05399
6	11.05	-13.05187

# Carry Out a Grid Search in R

```
## Plot log-likelihood vs. b0  
> ggplot(data = gridSearch, aes(x = b0, y = ll)) +  
  geom_line() +  
  theme_bw() +  
  xlab(expression(beta[0])) +  
  ylab("Log-Likelihood")
```

```
## Get the row with the largest log-likelihood  
> gridSearch %>% slice(which.max(ll))
```

Source: local data frame [1 x 2]

	b0	ll
	(dbl)	(dbl)
1	11.68	-12.9853

# Alternative: Use Calculus

- ★ The derivative of the log-likelihood function can be analytically computed and set equal to zero. Solving for  $\beta_0$  will give the value that maximizes the deviance
- ★ Avoids exhaustive search
- ★ In more complex models (e.g., GLMs) estimation is not so straightforward.
  - ✓ Exhaustive search methods are required (numerical analysis)
- ★ Once the log-likelihood function becomes more complex, the multidimensional form of the tangent line must be discovered

$$\frac{\partial}{\partial \beta_0} = -\frac{N}{2} \times \ln(2\pi\sigma_\epsilon^2) - \frac{1}{2\sigma_\epsilon^2} \times \sum (y_i - \beta_0)^2$$

Compute the partial derivative with respect to  $\beta_0$

Do you remember the product rule? What about the chain rule?


$$0 = -2[(y_1 - \beta_0) - (y_2 - \beta_0) - \dots - (y_n - \beta_0)]$$

Set this equal to 0 and solve the equation for  $\beta_0$  (find the roots of the function)


$$\beta_0 = \frac{\sum y_i}{n}$$

Sample mean is the estimate that maximizes the log-likelihood...in other words, the sample mean is the **maximum likelihood** estimate.

The calculus is essentially computing the  $\beta_0$  estimate that has a tangent line to the log-likelihood function with a slope of zero.




Log-Likelihood function



Likelihood function

Since the logarithm is a monotonic function, maximizing the log-likelihood function is equivalent to maximizing the likelihood function

## Example 2: Regression (Estimate Intercept, Slope, and Error Variance)



Given the data and the model, what is the most likely value of the unknown intercept parameter, slope parameter and error variance parameter?

**Data:** The data are given

**Model:** Regression model

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

Embed the model



$$\ell(\beta_0, \beta_1, \sigma_\epsilon^2; \mathbf{y}) = -\frac{5}{2} \times \ln(2 \times \pi \times \sigma_\epsilon^2) - \frac{1}{2 \times \sigma_\epsilon^2} \times \sum (y_i - \beta_0 - \beta_1(x_i))^2$$

$$\ell(\beta_0, \beta_1, \sigma_\epsilon^2; \mathbf{y}) = -\frac{5}{2} \times \ln(2 \times \pi \times \sigma_\epsilon^2) - \frac{1}{2 \times \sigma_\epsilon^2} \times [(12 - \beta_0 - \beta_1(32))^2 + (8 - \beta_0 - \beta_1(33))^2 + \dots + (8.5 - \beta_0 - \beta_1(26))^2]$$

Try alternative values for the three different parameters. Choose the set that gives the smallest deviance.

Candidate values for  $\beta_0$ ,  $\beta_1$ , and  $\sigma^2$  are considered simultaneously

$\beta_0$	$\beta_1$	$\sigma^2$	log-lik
-4.00	0	15	$d_1$
-3.99	0	15	$d_2$
-3.98	0	15	$d_3$
-3.97	0	15	$d_4$
$\vdots$	$\vdots$	15	$d_5$
-3.92	0.5	15	$d_6$
$\vdots$	$\vdots$	$\vdots$	$\vdots$

With three unknown parameters, we need to minimize a four-dimensional function.

Note that the search grid is **bigger by many magnitudes**.

Consider the following candidate values:

- $\beta_0 = [-4, -3]$  by 0.01 (100 candidate values)
- $\beta_1 = [0, 2]$  by 0.01 (200 candidate values)
- $\sigma^2 = [10, 15]$  by 0.01 (500 candidate values)

The search grid consists of  $100 \times 200 \times 500 = 10,000,000$  possible sets of candidate values!

**Computational time for me was about 14 minutes!**

And that is under a very specified range of values for each parameter and not much precision (hundredths).

Imagine if you needed the parameters estimated to 3 or 4 decimal places and had no clue what the values were?



# Log-Likelihood Function

```
## Create a function to compute the deviance
> LL2 = function(b0, b1, sigma2) {
  -5/2 * log(2 * pi * sigma2) - 1/(2 * sigma2) * (
    (12    - b0 - b1 * 32)^2 +
    (8     - b0 - b1 * 33)^2 +
    (16.26 - b0 - b1 * 32)^2 +
    (13.65 - b0 - b1 * 33)^2 +
    (8.5   - b0 - b1 * 26)^2
  )
}
```

```
## Try function
> LL2(b0 = 0, b1 = 0, sigma2 = 15)
```

```
[1] -35.73015
```

# Carry Out a Grid Search in R

```
## Generate candidate values for b0
> candidates = expand.grid(
  b0      = seq(from = -4, to = -3, by = 0.01)
  b0      = seq(from =  0, to =  1, by = 0.01)
  sigma2 = seq(from =  0, to = 20, by = 0.01)
)

## Generate the log-likelihood values and store the results
> gridSearch = candidates %>%
  rowwise() %>%
  mutate(l1 = dev(b0, b1, sigma2))

## Find the row with the smallest deviance
> gridSearch %>% slice(which.max(l1))
```

	b0	b1	sigma2	l1
	(dbl)	(dbl)	(dbl)	(dbl)
1	-3.92	0.5	7.99	-12.28932

```
## Examine search grid
## Note: Getting NAs does not necessarily mean your function failed...

> head(gridSearch)
```

	b0	b1	sigma2	ll
	(dbl)	(dbl)	(dbl)	(dbl)
1	-4.00	0	0	NA
2	-3.99	0	0	NA
3	-3.98	0	0	NA
4	-3.97	0	0	NA
5	-3.96	0	0	NA
6	-3.95	0	0	NA

Using calculus, we would minimize the function, by finding the roots for the three partial derivatives for the function setting each equal to zero, and solving for each parameter. In this case, using calculus saves us a lot of time.

It is convenient to use a computer to estimate the parameters and determine values of the log-likelihood.

```
## Create data
> x = c(32, 33, 32, 33, 26)
> y = c(12, 8, 16.26, 13.65, 8.5)

## Fit linear model
> lm.a = lm(y ~ x)

## Compute log-likelihood
> logLik(lm.a)

'log Lik.' -12.28932 (df=3)
```

# Deviance


Some statisticians prefer to work with the **deviance**.

$$-2 \times \ell(\boldsymbol{\beta}; \mathbf{y}) = -2 \times \left( -\frac{N}{2} \times \ln(2\pi\sigma_\epsilon^2) - \frac{1}{2\sigma_\epsilon^2} \times \sum (y_i - \beta_0)^2 \right)$$


The deviance function is the log-likelihood multiplied by  $-2$

$$\mathcal{D}(\boldsymbol{\beta}; \mathbf{y}) = N \times \ln(2\pi\sigma_\epsilon^2) + \frac{1}{\sigma_\epsilon^2} \times \sum (y_i - \beta_0)^2$$

Deviance function



Log-Likelihood function



Deviance function

Maximizing the log-likelihood function is equivalent to minimizing the deviance function

```
## Compute deviance  
> -2 * logLik(lm.a)[1]
```

```
24.57864
```

If we computed the deviance we would look for the MINIMUM value.


# Iterative Methods of Estimation

# Exhaustive Search vs. Numerical Methods

- ★ The exhaustive searches carried out in the examples were convenient
  - ✓ Ranges were used where the ML estimates were known to reside
  - ✓ In practice, these are not known
  - ✓ Increments are usually finer than 0.01 in practice
- ★ Computing time and memory becomes a valuable resource
- ★ Numerical methods, based on calculus are usually employed
  - ✓ Newton-Raphson method is most common algorithm
  - ✓ Deviance function assumed to be smooth and continuous with only one minimum (regularity assumption)
- ★ Functions such as `glm()` and `lmer()` generally combine both methods
  - ✓ Numerical methods are used to get in the neighborhood of the minimum deviance
  - ✓ Once more limited search space has been defined an iterative method can hone in on the minimum deviance until it is "good enough"



# Gradient Descent



# Gradient Descent

Iterative optimization algorithm to minimize a function,  $f(\beta)$

Begin with an initial guess for the values of the parameters

$$\beta_0^{(0)}, \beta_1^{(0)}, \dots, \beta_k^{(0)}$$

Compute the derivative of the function at the parameter value(s) and multiply by some small constant amount (called the learning rate),  $\alpha$


$$\alpha \left( \frac{\partial}{\partial \beta_j} f(\beta) \right)$$

Update the values of the parameters by

$$\beta_j^{(1)} = \beta_j^{(0)} - \alpha \left( \frac{\partial}{\partial \beta_j} f(\beta) \right)$$

Repeat using the new estimated solution


Stop when some criteria for convergence has been met




$$f'(x = -2) = -4 \quad f'(x = 2) = 4 \quad f'(x = 3) = 6$$

**Key:** Notice if derivative is negative the function slopes downward to the right; if derivative is positive the function slopes downward to the left; size of derivative indicates steepness of line

Gradient descent makes use of this to update the parameters



The minimum of the function is clearly at  $x = 0$ .



Start by making an initial guess of the parameter(s)...in this case  $x$

$$x^{(0)} = 4$$

Pick a learning rate (something small)

$\alpha = 0.1$  The value of the learning rate determines how fast the search converges.

Update the parameter,  $x$ , according to

$$\beta_j^{(1)} = \beta_j^{(0)} - \alpha \left( \frac{\partial}{\partial \beta_j} f(\beta) \right)$$

$$\frac{\partial}{\partial \beta_j} f(\beta) = 2x = 2(4) = 8 \quad \text{Evaluate at } x^{(0)}$$

$$x^{(1)} = 4 - 0.1(8) = 3.2$$

$x^{(1)}$  is closer to the minimum than  $x^{(0)}$

Repeat

$$x^{(2)} = 3.2 - 0.1(2 \cdot 3.2) = 2.56$$

$x^{(2)}$  is closer to the minimum than  $x^{(1)}$

Repeat until the decrease is less than some value

Here is an example of some output from an `lmer()` fitted model.

0:	2129.0419:
1:	2067.3454:
2:	2057.7418:
3:	2050.2435:
4:	2048.7592:
5:	2048.6277:
6:	2048.4524:
7:	2048.3884:
8:	2048.3731:
9:	2048.3708:
10:	2048.3705:
11:	2048.3705:
12:	2048.3705:
13:	2048.3705:
14:	2048.3705:

The first column indexes the iteration. The second gives the value of the deviance. Here the iterative estimation is done via penalized iteratively re-weighted least squares (PIRLS) rather than gradient descent.

# Problems with the Estimation

- ★ Two potential problems for iterative procedures
  - ✓ Possible for one of the  $\beta_k$  to **tend to infinity** (leads to non-convergence)
    - Generally means the model has been poorly specified (e.g., too many interactions, poor predictors, linear separation in a predictor, etc.)
    - Could also be due to data sparseness
    - Often the parameter can be fixed to a value to allow the other parameters to be estimated (the fixed parameter is then ignored in the output)
  - ✓ Under particular conditions it is possible for the **estimates to iterate in such a way that they never converge** (but don't tend toward infinity)
    - The likelihood (or log-likelihood) should be increasing at each iteration
    - If not, it is a sign that the true root has converged on a local maximum or will not converge
    - Step-halving functions can be applied where half the distance to the prior and current estimates is evaluated. If the likelihood is lower, the step-halving function is re-applied. If the likelihood increases, sub-iterations continue until there is no change in the likelihood

**Possible solution: Try different starting values**

**Possible solution: Fix some of the parameter values**